

Exascale Monte Carlo R&D

Summer 2012

Larry Cox, Ryan Marcus

Overview

- **Exascale computing**
 - Different technologies
 - Getting there
- **High-performance proof-of-concept: MCMini**
 - Monte Carlo neutron transport
 - Features
 - Results
- **OpenCL toolkit: Oatmeal**
 - Purpose
 - Features

Exascale Computing

- **558 times faster than Roadrunner**
- **Long way off**
 - Intel predicts 2018
- **Lots of different technologies**
- **Some commonalities between all technologies**
 - Host / Device paradigm
 - Parallel

Exascale Computing: Technologies

■ NVidia CUDA

- The original GPGPU language
- Proprietary: tied to NVidia devices

■ Khronos OpenCL

- A multi-platform massively parallel standard
- Runs on all current HPC hardware

■ Intel MIC

- 4 x86 CPUs packed onto a PCI card
- Runs current code (Fortran, C) with little modification

■ Advantages and disadvantages to each



Exascale Computing: Technologies

Quality	CUDA	OpenCL	MIC
Runs everywhere	✗	✓	✗
Open standard	✗	✓	✗
Pre-existing code	✗	✗	✓
Large user base	✓	✓	✗
CPU support	✗	✓	✓
Quality	CUDA	OpenCL	MIC
Backer	NVidia	Apple, AMD, Intel	Intel
API Support	Good	Moderate	Full

Exascale Computing: Technologies

- **Hardware-agnostic code is very important**
 - Write once, run anywhere
 - Changing landscape of HPC means compatibility is important
- **Running current code is very attractive, BUT it won't scale to the same level**
 - Memory constraints
 - Heavy instruction set
- **Open standard is key: code can't depend on one company**
- **MCMini's HPC technology: OpenCL**

MCMini


- **A Monte Carlo “mini-app”**
- **Proof of concept**
- **Performance capable and performance driven**
 - Written in C
- **Basic, Newtonian physics**
 - Algorithmically similar to “real” physics
- **Support for three reaction types: scatter, fission, and absorption**
- **Reads LNK3DNT meshes (MCNP geometry compatible)**
- **Multiple OpenCL devices (CPU, GPU)**
- **Multiple nodes**

MCMini: Capabilities


- MCMini can handle a lot of particles on really big geometries
- Trace and attenuation code is result-similar to MCNP
- Very capable (4th generation particles calculated)

Nodes	GPUs	Mesh cells	Particles	Time (m)
1	2	1000 ³	10 ⁵	0.5
2	4	2000 ³	10 ⁵	1.02
4	8	3000 ³	10 ⁵	2.8
8	16	10000 ³	10 ⁵	6.0


MCMini Scaling: CPU vs GPU (Cypress)



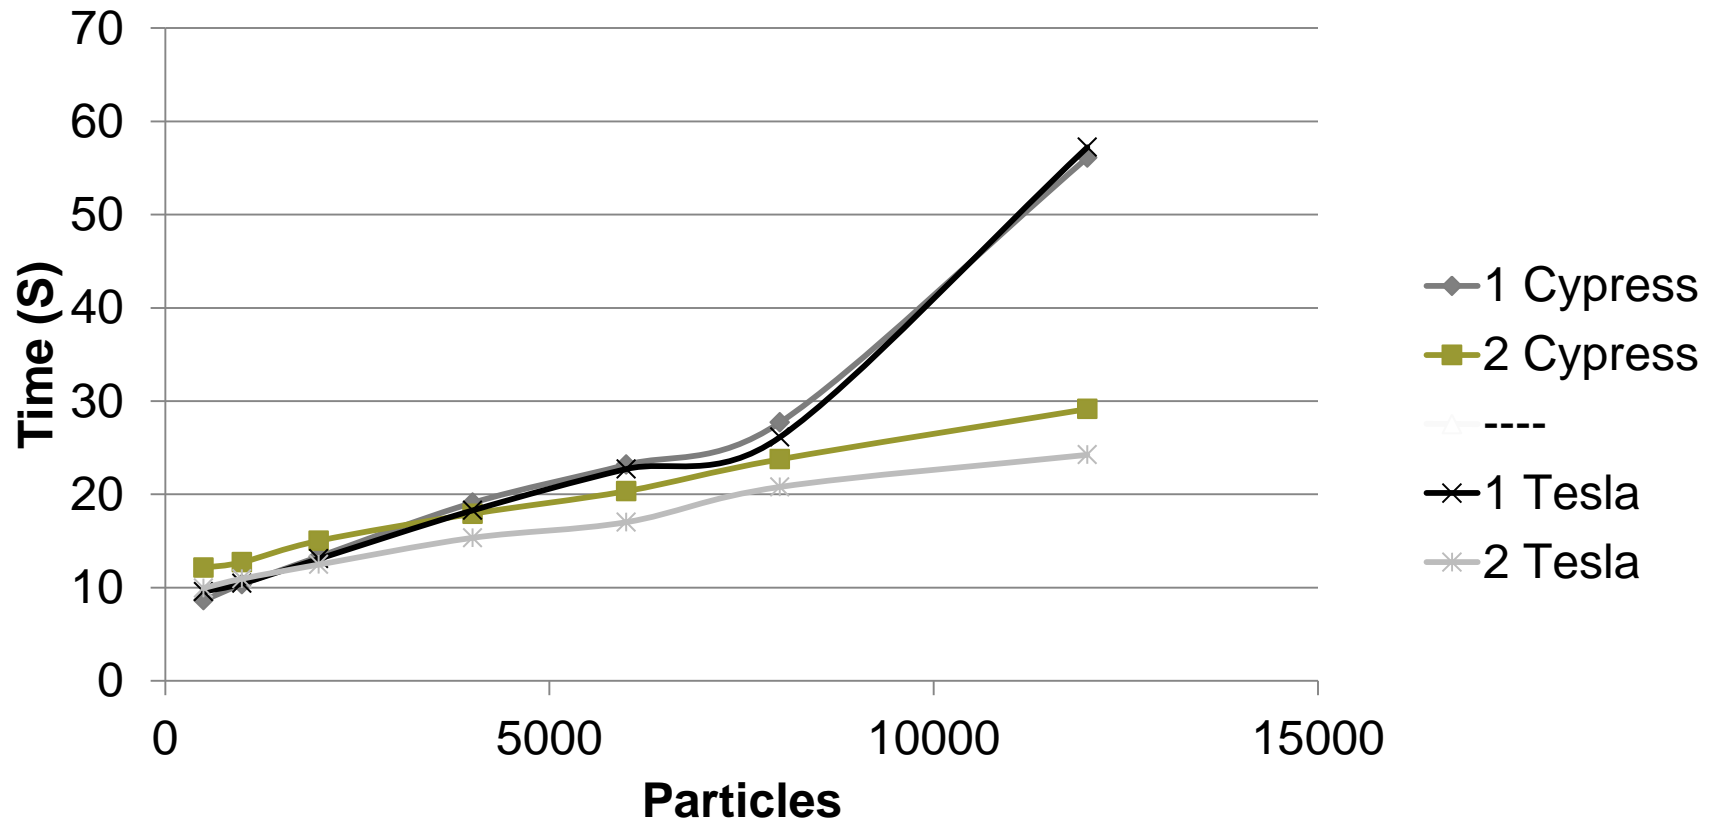
MCMMini Scaling: 1 GPU vs 2 GPUs (Cypress)



MCMini Scaling: Nodes



MCMini Scaling: Cypress vs Tesla (C2070)



MCMini / MCNP comparison

- **Neutron flux calculations**
 - “Comparable”
 - Similarly scaling
- **Time comparison is inherently unfair**
 - MCNP is doing far more than MCMini
- **Geometries that take a long time in MCNP generally requires minutes or seconds in MCMini**

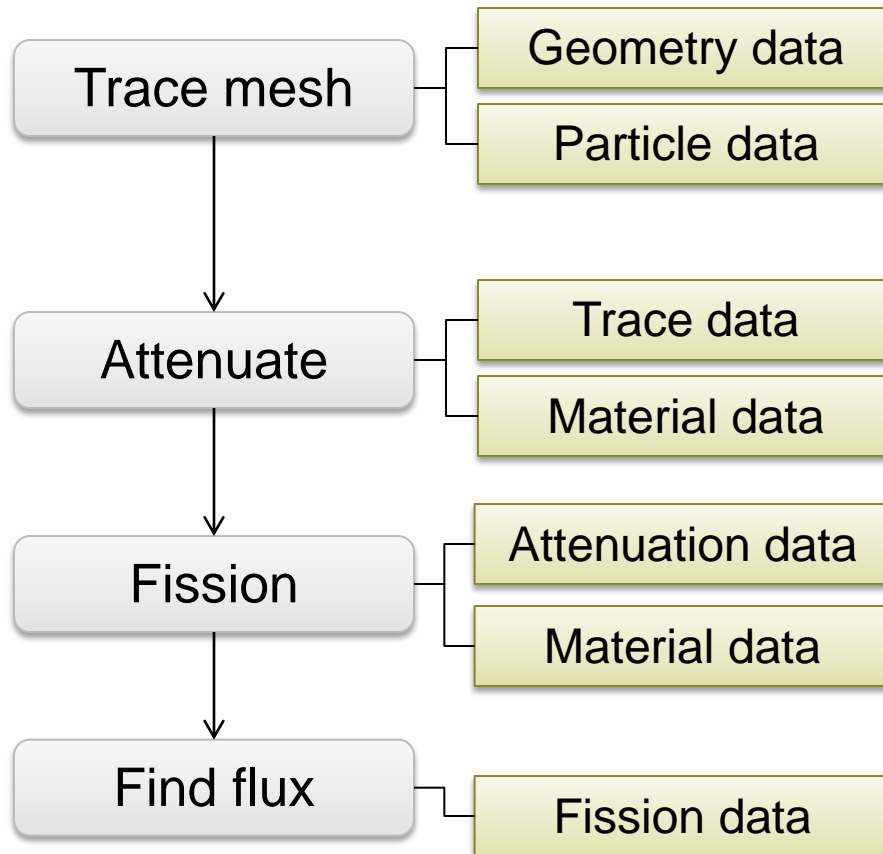
MCMini: Results

- **OpenCL is a viable HPC tool**
 - Scaling factors
 - General performance
- **OpenCL code can be easily designed to run on both CPUs and GPUs**
- **Monte Carlo style tasks are very well fitted to OpenCL and GPGPU computing**
- **Host/device paradigm introduces some interesting complications**

Oatmeal: Purpose

- **Traditionally, you either have enough RAM, or you don't.**
- **With GPGPUs**
 - Large host memory pool (500GB on some Darwin nodes)
 - Limited device memory pool (1 – 5 GB)
 - Device can't do computations on data not in device memory
- **Memory can be moved around between the host and the device**
 - Costly
- **OpenCL code is broken down into kernels**
 - Minimal task unit: the smallest logical division of a task

Oatmeal: Purpose



- **Steps require different data**
 - Data required by a step is that step's "context"
- **Context switching can be expensive**
 - Transferring data takes time
- **We want to minimize data transfer**
 - But we can't run out of memory
- **We could maintain a minimum context**
- **Or we can calculate an optimal context**

Oatmeal: Features

- Oatmeal is the OpenCL AuTomatic Memory Allocation Library
- A C++ framework for calculating optimal context switches
- Takes in a dynamic graph, executes the program
- Can run kernels over multiple devices and automatically reduce data
- Can run host-based tasks in parallel
- Can be used to profile memory usage and identify hotspots
 - Both memory and computational
- Can determine if bandwidth costs exceed computational benefit
 - And then run code on the CPU

Conclusions

- **Despite driver issues, OpenCL seems like a good, hardware agnostic tool**
- **MCMMini demonstrates the possibility for GPGPU-based Monte Carlo methods**
 - Shows great scaling for HPC application
 - Algorithmic equivalence
- **Oatmeal provides a flexible framework to aid in the development of scientific OpenCL codes**